
Par_incr: A library for incremental computation with support for parallelism

Dipesh Kafle | Mentor: Vesa Karvonen

What is incremental computation?

What is incremental computation?

- Software feature which, whenever a piece of data changes, attempts to save time by only recomputing those outputs which depend on the changed data.

What is incremental computation?

- Software feature which, whenever a piece of data changes, attempts to save time by only recomputing those outputs which depend on the changed data.
- Example: Excel, GUIs (e.g. React), Build systems, etc.

What is incremental computation?

- Software feature which, whenever a piece of data changes, attempts to save time by only recomputing those outputs which depend on the changed data.
- Example: Excel, GUIs (e.g. React), Build systems, etc.
- A bit more detailed description at [Introducing Incremental](#)

- Based on [Efficient Parallel Self-Adjusting Computation](#) (sort of a follow up paper to [Adaptive Functional Programming](#))

- Based on [Efficient Parallel Self-Adjusting Computation](#) (sort of a follow up paper to [Adaptive Functional Programming](#))

Existing similar libraries

- [current_incr](#)
- [Jane Street's incremental](#)

- Based on [Efficient Parallel Self-Adjusting Computation](#) (sort of a follow up paper to [Adaptive Functional Programming](#))

Existing similar libraries

- [current_incr](#)
- [Jane Street's incremental](#)

So what's different?

- Based on [Efficient Parallel Self-Adjusting Computation](#) (sort of a follow up paper to [Adaptive Functional Programming](#))

Existing similar libraries

- [current_incr](#)
- [Jane Street's incremental](#)

So what's different?

- You can leverage freshly introduced parallelism in `Par_incr`

How do we use the library?

How do we use the library?

We need to know about a few types before that

How do we use the library?

We need to know about a few types before that

- `'a Par_incr.Var.t`: Handle to the input of some computation.
Changing this should correctly propagate changes

How do we use the library?

We need to know about a few types before that

- `'a Par_incr.Var.t`: Handle to the input of some computation.
Changing this should correctly propagate changes
- `'a Par_incr.t`: This is different from `Var.t`. You call `Var.watch` on a `Var.t` to get `Par_incr.t`. The library provides combinators that operate on this to make larger computations.

How do we use the library?

We need to know about a few types before that

- `'a Par_incr.Var.t`: Handle to the input of some computation.
Changing this should correctly propagate changes
- `'a Par_incr.t`: This is different from `Var.t`. You call `Var.watch` on a `Var.t` to get `Par_incr.t`. The library provides combinators that operate on this to make larger computations.
- `'a Par_incr.computation`: When you run `'a Par_incr.t`, you'll obtain `'a computation`. This internally stores everything required to propagate changes.

How do we use the library?(continued)

```
1  module Cutoff : sig
2      type 'a t =
3          | Never
4          | Always
5          | Phys_equal
6          | Eq of ('a -> 'a -> bool)
7          | F of (oldval:'a -> newval:'a -> bool)
8
9      val attach : 'a t -> 'a incremental -> 'a incremental
10 end
```


How do we use the library?(continued)

```
1  module Var : sig
2    type 'a t
3    val create : ?cutoff:'a Cutoff.t -> ?to_s:('a -> string) -> 'a -> 'a t
4    val set : 'a t -> 'a -> unit
5    val value : 'a t -> 'a
6    val watch : 'a t -> 'a incremental
7    module Syntax : sig
8      val ( := ) : 'a t -> 'a -> unit
9      val ( ! ) : 'a t -> 'a
10   end
11 end
```

How do we use the library? (continued)

```
1  type 'a t
2  type 'a computation
3  type executor = {
4    run : 'a. (unit -> 'a) -> 'a;
5    par_do : 'a 'b. (unit -> 'a) -> (unit -> 'b) -> 'a * 'b;
6  }
7  module Cutoff: sig ... end
8  module Var : sig ... end
9  val return : 'a -> 'a t
10 val map : ?cutoff:'b Cutoff.t -> fn:('a -> 'b) -> 'a t -> 'b t
11 val map2 :
12   ?cutoff:'c Cutoff.t ->
13   ?mode:[`Par | `Seq] -> fn:('a -> 'b -> 'c) -> 'a t -> 'b t -> 'c t
14 val combine : 'a t -> 'b t -> ('a * 'b) t
15 val bind : fn:('a -> 'b t) -> 'a t -> 'b t
16 val par : left:'a t -> right:'b t -> ('a * 'b) t
17 val delay : (unit -> 'a t) -> 'a t
18 val value : 'a computation -> 'a
19 val run : executor:executor -> 'a t -> 'a computation
20 val propagate : 'a computation -> unit
21 val destroy_comp : 'a computation -> unit
```

How do we use the library? (continued)

How do we use the library? (continued)

- Define `Var.t` with some value.

How do we use the library? (continued)

- Define `Var.t` with some value.
- Perform `Var.watch` operation on `Var.t` and change it to `incremental(Par_incr.t)`.

How do we use the library? (continued)

- Define `Var.t` with some value.
- Perform `Var.watch` operation on `Var.t` and change it to `incremental(Par_incr.t)`.
- Use different combinators provided by the library on the `incrementals` and make even bigger `incrementals`.

How do we use the library? (continued)

- Define `Var.t` with some value.
- Perform `Var.watch` operation on `Var.t` and change it to `incremental(Par_incr.t)`.
- Use different combinators provided by the library on the `incrementals` and make even bigger `incrementals`.
- Obtain value of a certain `incremental` by running it.

How do we use the library? (continued)

- Define `Var.t` with some value.
- Perform `Var.watch` operation on `Var.t` and change it to `incremental(Par_incr.t)`.
- Use different combinators provided by the library on the `incrementals` and make even bigger `incrementals`.
- Obtain value of a certain `incremental` by running it.
- Running an `'a incremental` returns a `'a computation`.

How do we use the library? (continued)

- Define `Var.t` with some value.
- Perform `Var.watch` operation on `Var.t` and change it to `incremental(Par_incr.t)`.
- Use different combinators provided by the library on the `incrementals` and make even bigger `incrementals`.
- Obtain value of a certain `incremental` by running it.
- Running an `'a incremental` returns a `'a computation`.
- When we change a `Var.t`(done with `Var.set` operation), it marks all dependent computations dirty.

How do we use the library? (continued)

- Define `Var.t` with some value.
- Perform `Var.watch` operation on `Var.t` and change it to `incremental(Par_incr.t)`.
- Use different combinators provided by the library on the `incrementals` and make even bigger `incrementals`.
- Obtain value of a certain `incremental` by running it.
- Running an `'a incremental` returns a `'a computation`.
- When we change a `Var.t`(done with `Var.set` operation), it marks all dependent computations dirty.
- Running `propagate` operation on a dirty `computation` updates its value efficiently.

How do we use the library? (continued)

- Define `Var.t` with some value.
- Perform `Var.watch` operation on `Var.t` and change it to `incremental(Par_incr.t)`.
- Use different combinators provided by the library on the `incrementals` and make even bigger `incrementals`.
- Obtain value of a certain `incremental` by running it.
- Running an `'a incremental` returns a `'a computation`.
- When we change a `Var.t` (done with `Var.set` operation), it marks all dependent computations dirty.
- Running `propagate` operation on a dirty `computation` updates its value efficiently.
- Destroy (with `destroy_comp` operation) `computation` when its no more required.

Sequential sum_range

```
1  let rec sum_range ~lo ~hi xs =
2    Par_incr.delay @@ fun () ->
3      if hi - lo <= 1 then begin
4        xs.(lo)
5      end
6    else
7      let mid = lo + ((hi - lo) asr 1) in
8      Debug.attach ~fn:Int.to_string
9        Par_incr.Syntax.(
10         let+ lhalf = sum_range ~lo
11           ↪ ~hi:mid xs
12         and+ rhalf = sum_range ~lo:mid
13           ↪ ~hi xs in
14         lhalf + rhalf)
```

Example

Sequential sum_range

```
1 let rec sum_range ~lo ~hi xs =
2   Par_incr.delay @@ fun () ->
3   if hi - lo <= 1 then begin
4     xs.(lo)
5   end
6   else
7     let mid = lo + ((hi - lo) asr 1) in
8     Debug.attach ~fn:Int.to_string
9     Par_incr.Syntax.(
10      let+ lhalf = sum_range ~lo
11        ↪ ~hi:mid xs
12      and+ rhalf = sum_range ~lo:mid
13        ↪ ~hi xs in
14      lhalf + rhalf)
```

Parallel sum_range

```
1 let rec sum_range_par ~lo ~hi xs =
2   Par_incr.delay @@ fun () ->
3   if hi - lo <= 1 then begin
4     xs.(lo)
5   end
6   else
7     Debug.attach ~fn:Int.to_string
8     Par_incr.Syntax.(
9      let mid = lo + ((hi - lo) asr
10        ↪ 1) in
11      let& lhalf = sum_range_par ~lo
12        ↪ ~hi:mid xs
13      and& rhalf = sum_range_par
14        ↪ ~lo:mid ~hi xs in
15      lhalf + rhalf)
```

What happens internally?

Say we have something like this,

```
1  let count = 3
2  let arr = Array.init count (( + ) 1) (* [1,2,3] *)
3  let var_arr = Array.map (Par_incr.Var.create ~to_s:Int.to_string) arr
4  let t_arr = Array.map Par_incr.Var.watch var_arr
5  let executor =
6      Par_incr.
7      {
8          run = (fun f -> f ());
9          par_do =
10             (fun l r ->
11                 let lres = l () in
12                 (lres, r ()));
13      }
14  let () =
15      let seq_comp = Par_incr.run ~executor (sum_range ~lo:0 ~hi:count t_arr) in
16      Par_incr.dump_tree "sum-range.d2" seq_comp;
17      Var.set var_arr.(1) 10;
18      Par_incr.dump_tree "sum-range-after-change.d2" seq_comp;
19      Par_incr.propagate seq_comp;
20      Par_incr.dump_tree "sum-range-after-prop.d2" seq_comp;
21      Par_incr.destroy_comp seq_comp
```

What happens internally?(continued)

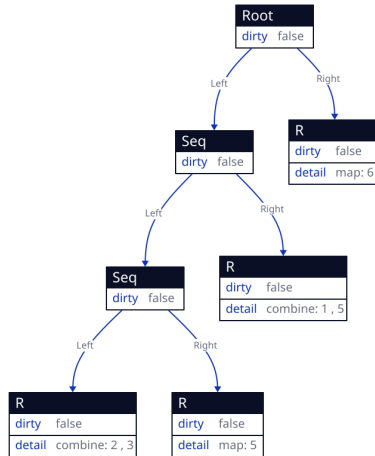


Figure 1: Internal representation of the computation

What happens internally?(continued)

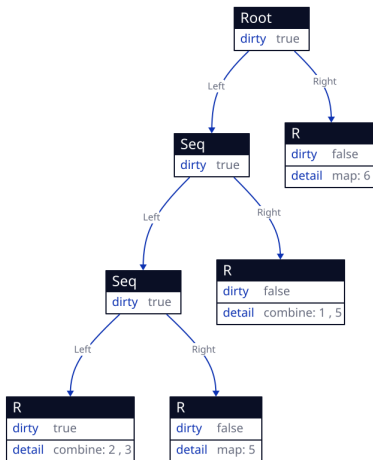


Figure 2: Computation after `var_arr.(1)` was changed

What happens internally?(continued)

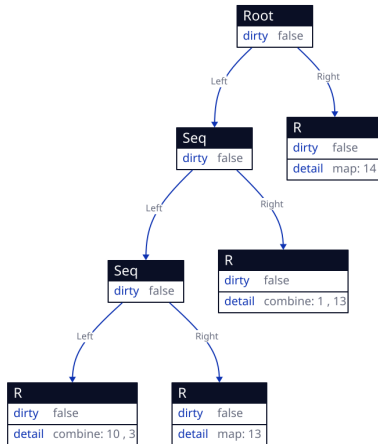


Figure 3: Computation after running propagate

What happens internally?(continue)

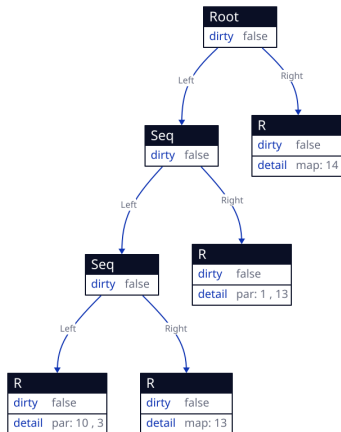


Figure 4: If we used `sum_range_par` instead

- Sequential performance within 2-4x of `Janestreet incremental` and pretty much always faster than `Current_incr`

- Sequential performance within 2-4x of `Janestreet incremental` and pretty much always faster than `Current_incr`
- Performance for parallelizable programs is almost always faster than `Janestreet incremental`. Sometimes by a lot and sometimes barely.

- Sequential performance within 2-4x of `Janestreet incremental` and pretty much always faster than `Current_incr`
- Performance for parallelizable programs is almost always faster than `Janestreet incremental`. Sometimes by a lot and sometimes barely.
- [Link to benchmarks](#)

- Diverting from the paper to a more functional style implementation

- Diverting from the paper to a more functional style implementation
- Improving performance

- Diverting from the paper to a more functional style implementation
- Improving performance
 - Having a standard documentation or a place to refer what optimizations kick in different cases would be super helpful.

- Diverting from the paper to a more functional style implementation
- Improving performance
 - Having a standard documentation or a place to refer what optimizations kick in different cases would be super helpful.
 - Compiler could inline stuff across module to some extent, but couldn't inline some larger functions even when I wanted to (with the `[@inlined]` annotation). Something like this should be possible, and it can help immensely in optimization experiments.

- Diverting from the paper to a more functional style implementation
- Improving performance
 - Having a standard documentation or a place to refer what optimizations kick in different cases would be super helpful.
 - Compiler could inline stuff across module to some extent, but couldn't inline some larger functions even when I wanted to (with the `[@inlined]` annotation). Something like this should be possible, and it can help immensely in optimization experiments.
- Profiling tools

- Library can probably be made faster, but it'll mostly require some major changes

- Library can probably be made faster, but it'll mostly require some major changes
- Implement some more programs using `Par_incr`

- Library can probably be made faster, but it'll mostly require some major changes
- Implement some more programs using **Par_incr**
- Implement the core of Differential Dataflow using **Par_incr**

THANK YOU
